

The Socratic Method: Cross-Resolution Supervision for Papyrus Surface Segmentation

ANONYMOUS AUTHOR(S)

SUBMISSION ID: TBD




Fig. 1. One raw checkpoint answers two opposing failure modes. Both rows begin with the identical CT crop and released M7 state-of-the-art baseline. The top row shows blind PHerc1447 de-blob behavior; the bottom row adds the projected fine teacher and shows longitudinal growth on locked PHerc0139. Every M7-XR panel uses the selected 8,192-sample student at $t = 0.45$, without blending.

Coarse-resolution papyrus surface segmentation fails in two opposing ways: faint sheets disappear, while nearby wraps merge into blobs. We present a teacher-student method for Herculaneum papyrus data in which an approximately $2.399 \mu\text{m}$ native-fine model interrogates the released $9.362 \mu\text{m}$ M7 segmenter. The resulting student remains an ordinary, single M7 residual-encoder nnU-Net at inference. During training, projected soft occupancy is combined with medial-crest recall, teacher-background separation, M7 function preservation, a global parameter trust region, and dynamic widest-path connectivity. A separate additive 2D curve fitter reconnects short, geometrically credible gaps after segmentation. Our completed duration study improves both held-out scrolls over released M7 at every measured interval through 8,192 samples, with a best calibrated macro-scroll Dice of 0.58293 at 7,168 samples. An independent registered-morphology and blind anti-blob review selects the raw 8,192-sample student at threshold 0.45. It passes the six-cube PHerc1447 anti-blob gate and matches teacher component count on 15 of 16 locked PHerc0139 slices; the remaining scalar exception and the difference between perceptual and overlap-based rankings are reported explicitly.

CCS Concepts: • **Computing methodologies** → **Neural networks**.

Additional Key Words and Phrases: volumetric segmentation, knowledge distillation, papyrus, topology-aware learning, curve completion

ACM Reference Format:

Anonymous Author(s). 2026. The Socratic Method: Cross-Resolution Supervision for Papyrus Surface Segmentation. *ACM Trans. Graph.* XX, X, Article XX (September 2026), 12 pages. <https://doi.org/10.1145/nnnnnnn.nnnnnnn>

1 Introduction

“I neither know nor think that I know.”

—Socrates, in Plato’s *Apology* 21d [Plato nd]

SOCRATES: I am wiser than this man; he fancies he knows something, although he knows nothing—

DARRYL, SOCRATES’ FRIEND: *fuck him up socrates*

—leon (@leyawn) [leon (@leyawn) 2015]

The objects behind this work are the Herculaneum papyri, a library of scrolls carbonized by the AD 79 eruption of Mount Vesuvius. Carbonization preserved the rolls but left them too brittle to open and the writing trapped inside tightly wound coils. Synchrotron micro-CT can record their internal structure without touching the object; recovering a page then requires locating the papyrus sheet in three dimensions and *virtually unwrapping* it to a plane [Seales et al. 2016; Vesuvius Challenge 2026b]. A scroll is one continuous sheet, metres long, visiting many adjacent turns. Consequently, two pieces separated by only a few CT voxels may be distant along the page, while a short gap in the volume may interrupt one continuous line of text.

Segmenting the papyrus surface is the first geometric commitment in a scroll-unwrapping pipeline. A missed surface interrupts a sheet that may otherwise be traceable over a long distance. An overgrown prediction is at least as dangerous: it can weld nearby turns into one component and make a locally plausible surface globally wrong. A useful segmentation method must therefore recover

faint structure and remove false connections at the same time. An overlap score alone does not adequately describe this tension.

The Vesuvius data provides a second, complementary source of information. A model operating on finer CT can resolve local evidence that is ambiguous after downsampling, while the deployed segmentation and geometry stack still consumes the coarser grid. We treat this mismatch as a cross-resolution dialogue. The fine model proposes soft occupancy and medial evidence; the released M7 model contributes a function that is already effective; explicit loss terms expose where those views are inconsistent. The student must reconcile them without moving far from M7.

This is the origin of our title. In the Socratic account, a speaker who professes little knowledge reveals inconsistencies in an expert's position and thereby points toward a better one. Our fine teacher is similarly not deployed as a superior replacement. It asks local questions of the expert coarse model. The answer is a revised M7 that stands alone at inference.

We present four contributions:

- a hash-pinned $2\mu\text{m}$ -to- $9\mu\text{m}$ teacher-supervision recipe that can be executed independently of our geometry repository;
- an objective that balances soft occupancy, de-blob separation, medial recall, dynamic connectivity, and two forms of M7 preservation;
- a duration-ladder protocol that retains each declared observation rather than assuming the longest run is best; and
- an independently measured, additive 2D curve fitter that reconnects only short gaps able to answer geometric and cross-plane objections.

The 8,192-sample duration experiment and the independent morphology review are complete. We report every committed interval, including the fact that the best-Dice interval is not the released interval. The selected artifact is the raw 8,192-sample student at threshold 0.45; its known component-count exception, undergrowth cases, and gate-proxy failures remain visible in the record.

2 Related Work

Segmentation. U-Net architectures and their self-configuring nnU-Net variants are widely used for dense volumetric segmentation [Isensee et al. 2021; Ronneberger et al. 2015]. Our deployed network follows the released M7 residual-encoder nnU-Net from the Vesuvius Challenge project exactly; we alter its supervision and optimization rather than adding an inference-time branch. This preserves compatibility with the existing coarse-volume pipeline.

Teacher supervision. Knowledge distillation transfers information from a teacher into a student by matching soft outputs or intermediate functions [Hinton et al. 2015]. Most such methods assume teacher and student predictions share a sampling grid. Our teacher observes a finer CT volume, so its occupancy and validity must be pulled back into the student's physical frame. The resulting targets represent partial volume rather than only a hard class.

Topology-aware Objectives. Skeleton and centerline losses preserve thin connected structures that can be under-represented by region overlap metrics [Shit et al. 2021]. The medial axis transform (MAT) instead describes a shape by centers and the radii of their

maximally inscribed disks or spheres. LSMAT casts this construction as a least-squares problem to make the resulting representation less brittle under noisy or perturbed boundaries [Rebain et al. 2019]. This center-radius factorization is especially useful for papyrus: motion along the medial set extends a sheet, whereas increasing radius thickens it toward a neighboring wrap. We use that distinction to derive thin training corridors and M7 contact pins. Unlike a conventional skeleton overlap loss, our final connectivity term does not prescribe every voxel of one centerline. It raises the bottleneck of whichever admissible medial path the student currently supports. The atlas is constructed only from fully owned training events and never from held-out data.

Curve completion. Classical thinning extracts one-pixel skeletons from binary images [Zhang and Suen 1984]. Curve completion can then be posed as endpoint pairing under geometric continuity. For wound papyrus, however, a locally smooth completion can cross from one wrap to the next. Our postprocess therefore combines 2D tangents with an umbilicus-aware radial test, neighboring-plane corridors, explicit anti-merger and crossing gates, and connection persistence across slices.

Vesuvius models and data. Our experiments build on the public Vesuvius Challenge CT data, the released M7 surface model, and a Villa native-fine teacher [Vesuvius Challenge 2026a,b; Villa Model Authors 2026]. The final paper will replace these artifact-level entries with the authoritative model and dataset citations associated with the public release.

3 Overview

Fig. 2 summarizes our pipeline. We first run a native-fine teacher in its approximately $2.399\mu\text{m}$ CT frame. Its probabilities, support, and medial crest are projected into the $9.362\mu\text{m}$ grid as soft training fields (Sec. 4). We initialize every candidate from the exact released M7 checkpoint and train all parameters under a small global trust region. The output is one ordinary M7 network. At inference we use only its raw class-1 probability; there is no teacher and no M7 probability blend.

After selecting a threshold, an optional C postprocess operates on each axial slice (Sec. 5). It proposes Bezier completions between skeleton endpoints, but paints a join only when local curve geometry, radial displacement, neighboring slices, and connection-track support agree. Segmentation training and curve completion are measured separately so an improvement in one stage cannot hide a failure in the other.

Two engineering details prevent accidental leakage or model substitution. First, the dynamic connectivity atlas is built only from boxes in the training manifest; held-out morphology is used only for evaluation. Second, the M7 initializer is accepted only if both its byte size and SHA-256 match the released artifact. A new candidate may resume its own interrupted optimizer state, but may not use an earlier student as learned initialization.




Fig. 2. Method overview with a registered crop. The released M7 appears both as the frozen initializer and as the state-of-the-art visual baseline. Training-only components (fine teacher, projection, medial and separation losses, and post-update trust projection) are separated from the deployed raw student. The optional 2D fitter is outside the learned model and is evaluated separately.

4 Cross-Resolution Supervision

4.1 Student Architecture and Data

Our student is a six-stage 3D residual-encoder nnU-Net with feature counts (32, 64, 128, 256, 320, 320), one CT input channel, and two output classes. It has 102,349,770 trainable parameters. The frozen training schedule contains 4,096 deterministic 192^3 patches from PHerc0139 and a single native-fine teacher record. We traverse this schedule twice for 8,192 samples. The small corpus reflects the time available before this progress submission, not the extent of the registered data. Our broader pair plan contains three additional non-prize training scrolls—PHerc0332, PHerc1667, and PHercParis4—with 62,500 candidate origins allocated to each. None of the four training scrolls is eligible for a Herculeaneum prize.

Our held-out validation manifest contains registered PHerc0814 and PHerc1451 patches. We evaluate every 1,024 samples and retain model-only checkpoints at 1,024, 2,048, 4,096, and 8,192 samples. Because the measured duration response is not monotonic (Sec. 6), we do not assume the final checkpoint is best.

4.2 Fine-Teacher Occupancy and Medial Geometry

The fine teacher provides two complementary targets in the student’s physical frame. The first is soft occupancy $q \in [0, 1]$, obtained by an anti-aliased pullback from the $2.399 \mu\text{m}$ grid. It estimates how much of a $9.362 \mu\text{m}$ voxel is occupied by papyrus. The second is a binary medial crest c , projected with an OR/max operator so that the existence of a fine-scale center survives downsampling. Occupancy and crest have independent validity masks.

This separation follows a concrete failure of occupancy-only supervision. A thin but real continuation commonly occupies only 25–45% of a coarse voxel. Soft cross-entropy is correctly calibrated when it drives such a voxel toward $p = 0.25$ – 0.45 , but the result may remain below the segmentation threshold. Reducing background

separation moved both the center and its shoulders upward, recovering length while reintroducing thick blobs. We therefore retain q as an amount-of-material target and use c as an independent sheet-existence target.

LSMAT-style center-radius representation. LSMAT represents a shape using medial centers and their maximally inscribed radii rather than boundary samples alone [Rebain et al. 2019]. We use that center-radius view while retaining the released Villa teacher’s center construction. For each native-fine axial slice T_z , the teacher skeletonizes the slice in 2D; the stacked centerlines receive one 3D morphological closing and are intersected with the original foreground. We pair every surviving center $u \in C_z$ with a physical 2D Euclidean-distance radius

$$r_z(u) = \min_{y \notin T_z} \|u - y\|_{\text{phys}}, \quad \hat{T}_z = \bigcup_{u \in C_z} B^\circ(u, r_z(u)). \quad (1)$$

The disks are open because the stored distance reaches the nearest background voxel center. Slice-wise radii are deliberate: pairing 2D centers with a 3D distance transform would mix disks and spheres and collapse the desired medial surface toward a curve. We audit reconstruction, component count, false-positive girth, and halo stability before using these centers for supervision.

The representation gives us two independent geometric coordinates. Moving along C_z extends the papyrus sheet tangentially; increasing r_z grows it radially toward another wrap. That distinction is the reason for introducing the LSMAT view: the connectivity loss should act along the first coordinate, while the separation loss and radius screens police the second.

From medial centers to pins. We construct events only on training slices with at least 95% valid teacher support. Components smaller than 12 voxels are removed. A slice qualifies when one thin teacher component H_e contacts at least two disconnected M7 components

$M_{e,i}$ and contains at least two teacher voxels absent from the one-voxel M7 neighborhood. Both teacher and M7 must have maximum in-plane EDT radius at most three voxels, and no more than 2% of their foreground may survive a two-voxel erosion; these checks reject blobs before they can define a positive connectivity target.

Let C_e be the projected medial centers inside H_e , and let D_r denote radius- r dilation. We choose the smallest $d \in \{0, 1, 2, 3\}$ for which

$$\Omega_e = H_e \cap D_d(C_e), \quad A_{e,i} = \Omega_e \cap D_2(M_{e,i}) \quad (2)$$

contains every nonempty contact set $A_{e,i}$ in one connected component. The sets $A_{e,i}$ are the *pins*: they are where distinct released-M7 fragments meet the LSMAT-style teacher-medial corridor, not arbitrary endpoints chosen from a student prediction. Corridor voxels within one voxel of M7, plus all pin voxels, form the free-anchor set F_e . Events touching their owning tile boundary are discarded, so every retained corridor is transformed and loaded as one unit.

4.3 Objective

Let p be the student's surface probability, q the projected fine-teacher occupancy, c its binary medial-crest field, $C = \{x : c_x = 1\}$ the corresponding valid crest set, and p_0 the released M7 probability. Our loss is

$$\begin{aligned} \mathcal{L} = & \mathcal{L}_{\text{CE}}(p, q) + 0.25\mathcal{L}_{\text{Dice}}(p, q) + \mathcal{L}_{\text{crest}}(p, c) \\ & + 2\mathcal{L}_{\text{sep}}(p, q) + 0.5\mathcal{L}_{\text{KL}}(p, p_0) + \mathcal{L}_{\text{pres}}(p, p_0) \\ & + 0.03125\mathcal{L}_{\text{conn}}(p). \end{aligned} \quad (3)$$

All terms are evaluated only where their corresponding validity fields permit.

Soft Occupancy and Crest Recall. Our cross-entropy and Dice targets retain fractional occupancy created by fine-to-coarse projection. For a nonempty valid crest C , the additional term is

$$\mathcal{L}_{\text{crest}} = 1 - \frac{\sum_{x \in C} p_x + 1}{|C| + 1}. \quad (4)$$

We average Eq. 4 over crest-bearing samples rather than over all voxels. The normalization prevents sparse medial evidence from disappearing into the background count, and the additive one is the smoothing constant used by the released teacher loss. Because only crest voxels receive its gradient, the term asks for longitudinal sheet evidence without defining a thick target.

De-Blob Separation. Thick predictions need explicit shrink pressure. An M7 blob can overlap the fine sheet and still score well against an anti-aliased occupancy target. Let V and V_c denote occupancy and crest validity. We form the thin positive seed

$$P = C \cup (\{q \geq 0.5\} \cap (V \setminus V_c)), \quad (5)$$

using the crest wherever it is known and falling back voxelwise to hard occupancy only where it is not. We dilate P by two voxels and retain only occupancy-valid locations where the teacher is confidently background:

$$S = (\text{dilate}_2(P) \setminus P) \cap V \cap \{q \leq 0.1\}. \quad (6)$$

The separation loss is the mean background cross-entropy

$$\mathcal{L}_{\text{sep}} = -\frac{1}{|S|} \sum_{x \in S} \log(1 - p_x). \quad (7)$$

We weight this term by 2 in Eq. 3. Because S is a local fence around supported medial structure rather than arbitrary background, the term attacks thickness and short cross-wrap welds while the soft occupancy and crest terms remain free to grow a faint sheet elsewhere.

M7 function and positive preservation. The teacher is locally more informative, not globally complete. We therefore freeze a copy of released M7 during training and use its probability p_0 in two differently masked terms. Let $H = C \cup (\{q \geq 0.5\} \cap V)$, $h_x = \mathbf{1}[x \in H]$, and $m_x = \mathbf{1}[p_{0,x} \geq 0.5]$. The KL mask is

$$\begin{aligned} K = & ((\bar{V}) \setminus D_2(H)) \\ & \cup \{x \in V : h_x = m_x \wedge (q_x \leq 0.1 \vee h_x = 1)\}. \end{aligned} \quad (8)$$

Thus unknown space is anchored except for a radius-two corridor around known teacher positives. In known space, M7 is anchored only where its hard decision agrees with a confident teacher decision; ambiguous partial occupancy is not misclassified as agreement on background. We use

$$\mathcal{L}_{\text{KL}} = \frac{1}{|K|} \sum_{x \in K} D_{\text{KL}}(\text{Ber}(p_{0,x}) \parallel \text{Ber}(p_x)). \quad (9)$$

KL preservation alone can still allow a correct M7 foreground fragment to shrink under the separation term. We therefore define $R = V \cap D_2(H) \cap \{p_0 \geq 0.5\}$ and add the positive-only loss

$$\mathcal{L}_{\text{pres}} = -\frac{1}{|R|} \sum_{x \in R} \log p_x. \quad (10)$$

The two terms have separate roles: Eq. 9 retains M7's foreground and background function away from justified edits, while Eq. 10 prevents de-blobbing from deleting M7-supported sheet near the teacher. Teacher disagreement remains free to correct M7.

Dynamic medial connectivity. Separation supplies the desired shrink pressure, but without a complementary topological term it can also sever the same faint longitudinal connections we want to retain. Our first LSMAT axial loss connected the same pins with a minimum-off-axis spanning tree, then raised the weakest 10% of voxels on that fixed route toward probability 0.2. This proved that a thin, center-biased route could be supervised without using held-out geometry, but it also froze one of many equally valid paths. The lowest weight only tied the incumbent; larger weights monotonically reduced ordinary validation quality, and the strongest weight also reduced anti-blob passes. The final loss therefore retains the audited LSMAT corridor and pins but makes the path existential.

For event e , define the path capacity

$$a_e(x) = \begin{cases} 1, & x \in F_e, \\ p_x, & x \in \Omega_e \setminus F_e, \\ 0, & x \notin \Omega_e. \end{cases} \quad (11)$$

Taking $A_{e,0}$ as a source pin, initialize $r_e^0(x) = \mathbf{1}[x \in A_{e,0}]$ and iterate

$$r_e^{t+1}(x) = \mathbf{1}[x \in \Omega_e] \max(r_e^t(x), \min(a_e(x), \max_{y \in \mathcal{N}_{26}(x)} r_e^t(y))). \quad (12)$$

This max-min recurrence computes the greatest bottleneck capacity among paths of at most t steps. After 96 steps, the event score and

loss are

$$b_e = \min_{i>0} \max_{x \in A_{e,i}} r_e^{96}(x),$$

$$\mathcal{L}_{\text{conn}} = \frac{1}{|E|} \sum_{e \in E} \max(0, \text{logit}(0.2) - \text{logit}(b_e)). \quad (13)$$

Each event receives one vote regardless of corridor size or pin count. Max/min back-propagation places pressure on the current best path's limiting capacity, not on every corridor voxel. Free M7 anchors have unit capacity and no direct gradient; no connectivity gradient exists outside Ω_e . The loss becomes exactly zero when every pin is reachable through some path with bottleneck at least 0.2.

The frozen atlas contains 149 fully owned events. Ninety-nine rows in the exact schedule encounter an event, 477 rows are eligible, and the longest audited event needs 44 of the allocated 96 propagation steps. Of these events, 135 have two pins, 13 have three, and one has four; 122 need no crest dilation and 27 need one voxel. The loss is evaluated only at the full-resolution output and has weight 0.03125. Together, the separation shell and existential path constraint express the two sides of de-blobbing: remove unsupported girth without deleting sheet existence.

Deep supervision. The occupancy, crest, separation, KL, and preservation terms are applied at the standard nnU-Net decoder outputs with normalized weights proportional to $1, 1/2, 1/4, \dots$ and zero weight on the coarsest output. Soft occupancy is resized trilinearly. The sparse crest is max-pooled, while its validity is true only if every contributing fine cell is valid. Connectivity is intentionally excluded from lower-resolution outputs so resampling cannot broaden its medial corridor.

4.4 Optimization and Trust Region

We use Adam with constant learning rate 2×10^{-5} , betas $(0.9, 0.999)$, $\epsilon = 10^{-8}$, and zero weight decay. The batch size is three with eight-step gradient accumulation and bfloat16 autocast. Seed 1203 fixes the sample order and augmentation sequence.

After each optimizer update, we project the complete parameter vector θ into a relative- L^2 ball around released parameters θ_0 :

$$\theta \leftarrow \theta_0 + \min\left(1, \frac{r \|\theta_0\|_2}{\|\theta - \theta_0\|_2}\right) (\theta - \theta_0), \quad r = 0.0027535422. \quad (14)$$

The projection is active on essentially every observed update. It is thus part of the effective method rather than an inactive guardrail.

5 Additive 2D Curve Fitting

Following re-training, we introduce a downstream postprocess that repairs short omissions in the thresholded predictions. It is additive only: accepted joins set foreground pixels, but never erase the model output. This restriction makes every intervention easy to localize and prevents a gap-repair stage from quietly becoming a second segmentation method.

For each z -slice, we apply Zhang–Suen thinning [Zhang and Suen 1984], remove short spurs, label the 8-connected skeleton and mask components, and extract curve endpoints with outward tangents. We then snap endpoints to the mask edge so distances describe the visible gap rather than skeleton retraction. Endpoints clipped by

the volume or missing data are retained for audit but may not be matched.

We score endpoint pairs using distance, facing, adjacent-plane evidence, and cross-slice track support. Before a pair enters deterministic greedy matching it must satisfy the following requirements:

- (1) it does not close the same skeleton or mask component;
- (2) its tangents face the gap and are sufficiently opposed;
- (3) its reach is at most 21 pixels, while the calibrated safe regime at six pixels or less requires less supporting evidence;
- (4) if an umbilicus is available, its radial displacement is at most three pixels;
- (5) its cubic Bezier has a bounded arc-to-chord ratio and does not approach a third component;
- (6) the path corridor does not cross another sheet on nearby slices; and
- (7) it neither crosses nor touches an already accepted join, and neither endpoint has taken a better match.

A first matching pass seeds endpoint and connection tracks. A second pass follows each new track through z , uses connection confidence, and retains only joins supported across at least three slices. We then paint a cubic Bezier disk stroke and record every added pixel.

6 Results

6.1 Evaluation Protocol

We calibrate thresholds from 0.25 through 0.60 in increments of 0.01 on registered PHerc0814 and PHerc1451 patches. Checkpoint ranking uses calibrated macro-scroll Dice with a minimum per-scroll gain guard against released M7. Promotion additionally requires a blinded PHerc1447 anti-blob review, two-sided morphology and coverage checks, and postprocessor false-join review. An earlier raw v29 candidate failed the PHerc1447 gate; we do not consider it a viable model.

6.2 Duration Study

Table 1 reports every committed validation interval through the completed 8,192-sample schedule. All intervals improve both held-out scrolls over released M7. The response is shallow and non-monotonic; the best ordinary validation result is the 7,168-sample interval at 0.58293 calibrated macro-scroll Dice. This behavior justifies retaining the full ladder and separating validation ranking from the morphology-based release decision.

The trust-region projection is active for 97.7% of optimizer updates in the first interval and 100% thereafter. The method therefore reaches its best observed validation score while remaining on the boundary of its allowed parameter neighborhood.

Every interval selects threshold 0.25, the lower edge of our ordinary Dice sweep. The optimum is censored, and the corresponding predictions inflate foreground on blind PHerc1447; 0.25 is therefore rejected. Independent registered morphology and anti-blob review selects the raw 8,192-sample student at $t = 0.45$.

At the selected point, the six PHerc1447 cubes contain 0.951 times the foreground of the v15 anti-blob comparison and retain 0.911 of that comparison's foreground. All aggregate and per-cube anti-blob constraints pass, with no interior-fraction or maximum-thickness



Fig. 3. Actual PHerc0139 training data at dynamic event 138. Projected soft occupancy, its sparse medial crest, and the radius-two separation shell encode center/radius geometry without rewarding girth. The lower panels show the M7 fragments, teacher-medial corridor, component-contact pins, and free anchors supplied to the max-min connectivity recurrence.




Fig. 4. Measured line-fitter behavior on the released M7 baseline. Left: representative M7 masks before and after additive accepted joins. Center: the same connection persists across neighboring z-slices. Right: registered fine-CT evidence distinguishes a supported connection from a required separation. The calibration accepted 481 of 772 proposals, added 12,697 pixels, reached 92.5% registered fine-CT precision, and produced no observed full-turn fusion. Cyan is released M7 and green is fitted addition.

regressions. On the 16 locked PHerc0139 slices, all anti-blob checks

pass, mean ASSD is 0.626 voxels, and 15 component counts match.

Table 1. Completed v31 duration ladder. “Cal.” is calibrated macro-scroll Dice; “ $t = .40$ ” uses a fixed threshold; “gain” is calibrated macro gain over the M7 initial model. All calibrated thresholds equal the 0.25 search boundary.

Samples	Cal.	$t = .40$	Gain	Min. gain
1,024	.57944	.57136	+0.01968	+0.01917
2,048	.57886	.57086	+0.01910	+0.01868
3,072	.58037	.57253	+0.02062	+0.02037
4,096	.58022	.57196	+0.02047	+0.02024
5,120	.57873	.57014	+0.01897	+0.01813
6,144	.58104	.57329	+0.02128	+0.02124
7,168	.58293	.57537	+0.02317	+0.02276
8,192	.58037	.57226	+0.02061	+0.02031

Rank 26 is the exception: the teacher truly has four components, while two top blobs and a touching pair of strokes yield five components in the student despite recovering the intended structures.

FLIP mean, which rewards literal teacher resemblance, narrowly prefers 4,096 samples at $t = 0.42$ (0.22836 versus 0.23331 for the release). This is not hidden: the selected point ranks 18 of 45 anti-blob-eligible settings by that statistic. Against the 2,048-sample checkpoint at fixed $t = 0.45$, however, 8,192 improves FLIP on 12 of 16 slices while reducing both teacher-missing pixels (4,283 to 4,270) and additions (1,201 to 1,155). Human review and blind de-blob behavior resolve the remaining choice in favor of the longer run.

6.3 Line-Fitter Evaluation

In a calibrated $4 \times 5 \times 5$ study, the line fitter accepted 481 joins and added 12,697 pixels. Registered high-resolution precision was 92.5% overall and 98.4% for joins no longer than six pixels. We observed no full-turn fusion and measured a 21% reduction in atlas overlap. A larger 21-cube census measured 98.3% weighted precision.

Adjacent high-resolution connectivity is useful confirmation but is not, by itself, proof that a join remains on one wrap. The umbilicus-aware radial gate is our primary cross-wrap safeguard. The release evaluation will therefore report false joins by gate, distance, and radial regime rather than only an aggregate precision.

6.4 Interpretation

The duration experiment establishes consistent positive ordinary held-out gains under a saturated M7 trust constraint. The independent gates establish a selected raw checkpoint and operating threshold, but not ground-truth accuracy: the PHerc0139 teacher and PHerc1447 v15 reference are both model-derived. The release record therefore preserves metric exceptions and human-review decisions alongside aggregate scores rather than collapsing them into one pass bit.

7 Limitations

Our fitting schedule is PHerc0139-only and may encode scroll-specific material or acquisition properties. Fine-teacher predictions are model-derived soft supervision rather than human ground truth. Two held-out scrolls provide a guard but do not establish broad robustness. The M7 trust projection is saturated, and its radius was selected under an earlier objective rather than the final loss in Eq. 3.

The ordinary Dice sweep also fails to bracket its optimum; threshold 0.45 is an independent morphology operating point, not the Dice optimum.

The line fitter has a different failure mode. A completion may be locally smooth and persistent across slices but still join the wrong wrap near the core. Radial, third-component, and cross-sheet gates reduce this risk but cannot prove global correctness. The postprocess must remain independently measured and its additions must remain recoverable during qualification.

Finally, exact replay requires large, separately licensed artifacts. The training code, recipes, hashes, and export scaffolding are public-repository material, but the approximately 31 GB validation corpus, teacher atlases, M7 weights, and student checkpoints require dedicated artifact hosting and resolved upstream terms.

8 Conclusion

We presented a cross-resolution supervision method for improving the released M7 papyrus-surface segmenter without changing its deployed architecture. A fine teacher exposes local inconsistencies through soft occupancy, medial evidence, and de-blob separation; M7 preservation and a global trust region keep the answer close to the incumbent function. The same design principle extends downstream: the 2D fitter may propose a connection, but radial, merger, cross-sheet, crossing, and persistence gates force it to answer a sequence of geometric objections before adding pixels.

The run improves both ordinary held-out scrolls through every measured interval and reaches its best observed macro-scroll Dice at 7,168 samples. Independent morphology and anti-blob evidence selects the raw 8,192-sample checkpoint at $t = 0.45$ instead. This separation between score optimization and geometric judgment is central to the method: a useful surface model must learn to grow longitudinal structure without buying that recall through thickness or cross-wrap mergers.




Fig. 5. Locked PHerc0139 comparisons at the selected operating point. Columns show the same CT crop, released M7, projected fine-teacher target, raw student, and a two-sided teacher comparison. Green denotes shared foreground, yellow teacher-only foreground, and magenta student-only foreground.




Fig. 6. Known edge cases at $t = 0.45$: a component-count proxy that disagrees with the visible structure, mild drift or thinning on locked sheets, and blind PHerc1447 views that remain undergrown or fragmented. Every row includes CT, released M7, the available model-derived reference, raw M7-XR, and a two-sided comparison. Yellow reference masks are not ground truth.

9 Qualitative Gallery

The following pages collect the clearest qualitative results from the selected raw checkpoint. They are presentation views, not additional samples or a hidden selection stage: every crop appears in the committed locked-16 or fixed PHerc1447 review corpus, and every student panel uses 8,192 samples at $t = 0.45$ without blending.



Fig. 7. Locked PHerc0139 growth gallery. Every registered crop is shown as CT, released M7, projected fine-teacher target, and raw M7-XR. These examples emphasize learned longitudinal regularity while keeping the released state of the art visible.



Fig. 8. Blind PHerc1447 anti-blob gallery. The released M7 foreground is cyan and the selected raw student's foreground is magenta, each on the identical CT slice. No fine teacher exists for these views.

References

Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. 2015. Distilling the Knowledge in a Neural Network. In *NIPS Deep Learning and Representation Learning Workshop*. arXiv:1503.02531

Fabian Isensee, Paul F. Jaeger, Simon A. A. Kohl, Jens Petersen, and Klaus H. Maier-Hein. 2021. nnU-Net: a self-configuring method for deep learning-based biomedical image segmentation. *Nature Methods* 18 (2021), 203–211. doi:10.1038/s41592-020-01008-z

leon (@leyawn). 2015. Socrates and Darryl. Twitter post. <https://twitter.com/leyawn/status/585859882869469184> Posted 2015-04-08.

Plato. n.d.. Apology, Section 21d. Perseus Digital Library. <https://www.perseus.tufts.edu/hopper/text?doc=Perseus%3Atext%3A1999.01.0170%3Atext%3DApol.%3Asection%3D21d> Accessed 2026-09-01.

Daniel Rebain, Baptiste Angles, Julien Valentin, Nicholas Vining, Jiju Peethambaran, Shahram Izadi, and Andrea Tagliasacchi. 2019. LSMAT: Least Squares Medial Axis Transform. *Computer Graphics Forum* 38, 6 (2019), 5–18. doi:10.1111/cgf.13599

Olaf Ronneberger, Philipp Fischer, and Thomas Brox. 2015. U-Net: Convolutional Networks for Biomedical Image Segmentation. In *Medical Image Computing and Computer-Assisted Intervention*. 234–241. doi:10.1007/978-3-319-24574-4_28

William Brent Seales, Clifford Seth Parker, Michael Segal, Emanuel Tov, Pnina Shor, and Yosef Porath. 2016. From Damage to Discovery via Virtual Unwrapping: Reading the Scroll from En-Gedi. *Science Advances* 2, 9 (2016), e1601247. doi:10.1126/sciadv.1601247

Suprosanna Shit, Johannes C. Paetzold, Anjany Sekuboyina, Andrey Zhylka, Ivan Ezhov, Alexander Unger, Jose P. W. Pluim, Ulrich Bauer, and Bjoern H. Menze. 2021. cLDice – A Novel Topology-Preserving Loss Function for Tubular Structure Segmentation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*. 16560–16569.

Vesuvius Challenge. 2026a. Released M7 Surface Segmentation Model. Model authors, release URL, and version to be finalized before submission.

Vesuvius Challenge. 2026b. Vesuvius Challenge Data and Tasks. Artifact citation and access date to be finalized before submission.

Villa Model Authors. 2026. Native-Fine Papyrus Surface Segmentation Teacher. Authoritative software and checkpoint citation to be finalized before submission.

T. Y. Zhang and C. Y. Suen. 1984. A Fast Parallel Algorithm for Thinning Digital Patterns. *Commun. ACM* 27, 3 (1984), 236–239. doi:10.1145/357994.358023